

Concurrency: Deadlock

Chapter 6

Organization

- ▶ **Principles of Deadlock**
- ▶ **Deadlock Prevention**
- ▶ **Deadlock Avoidance**
- ▶ **Deadlock Detection**



Principles of Deadlock

Deadlock: The Definition


- ▶ **Definition**

- ▶ *Permanent blocking of a set of processes that either compete for system resources or communicate with each other*
- ▶ *A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.*


- ▶ **No efficient solution!!**



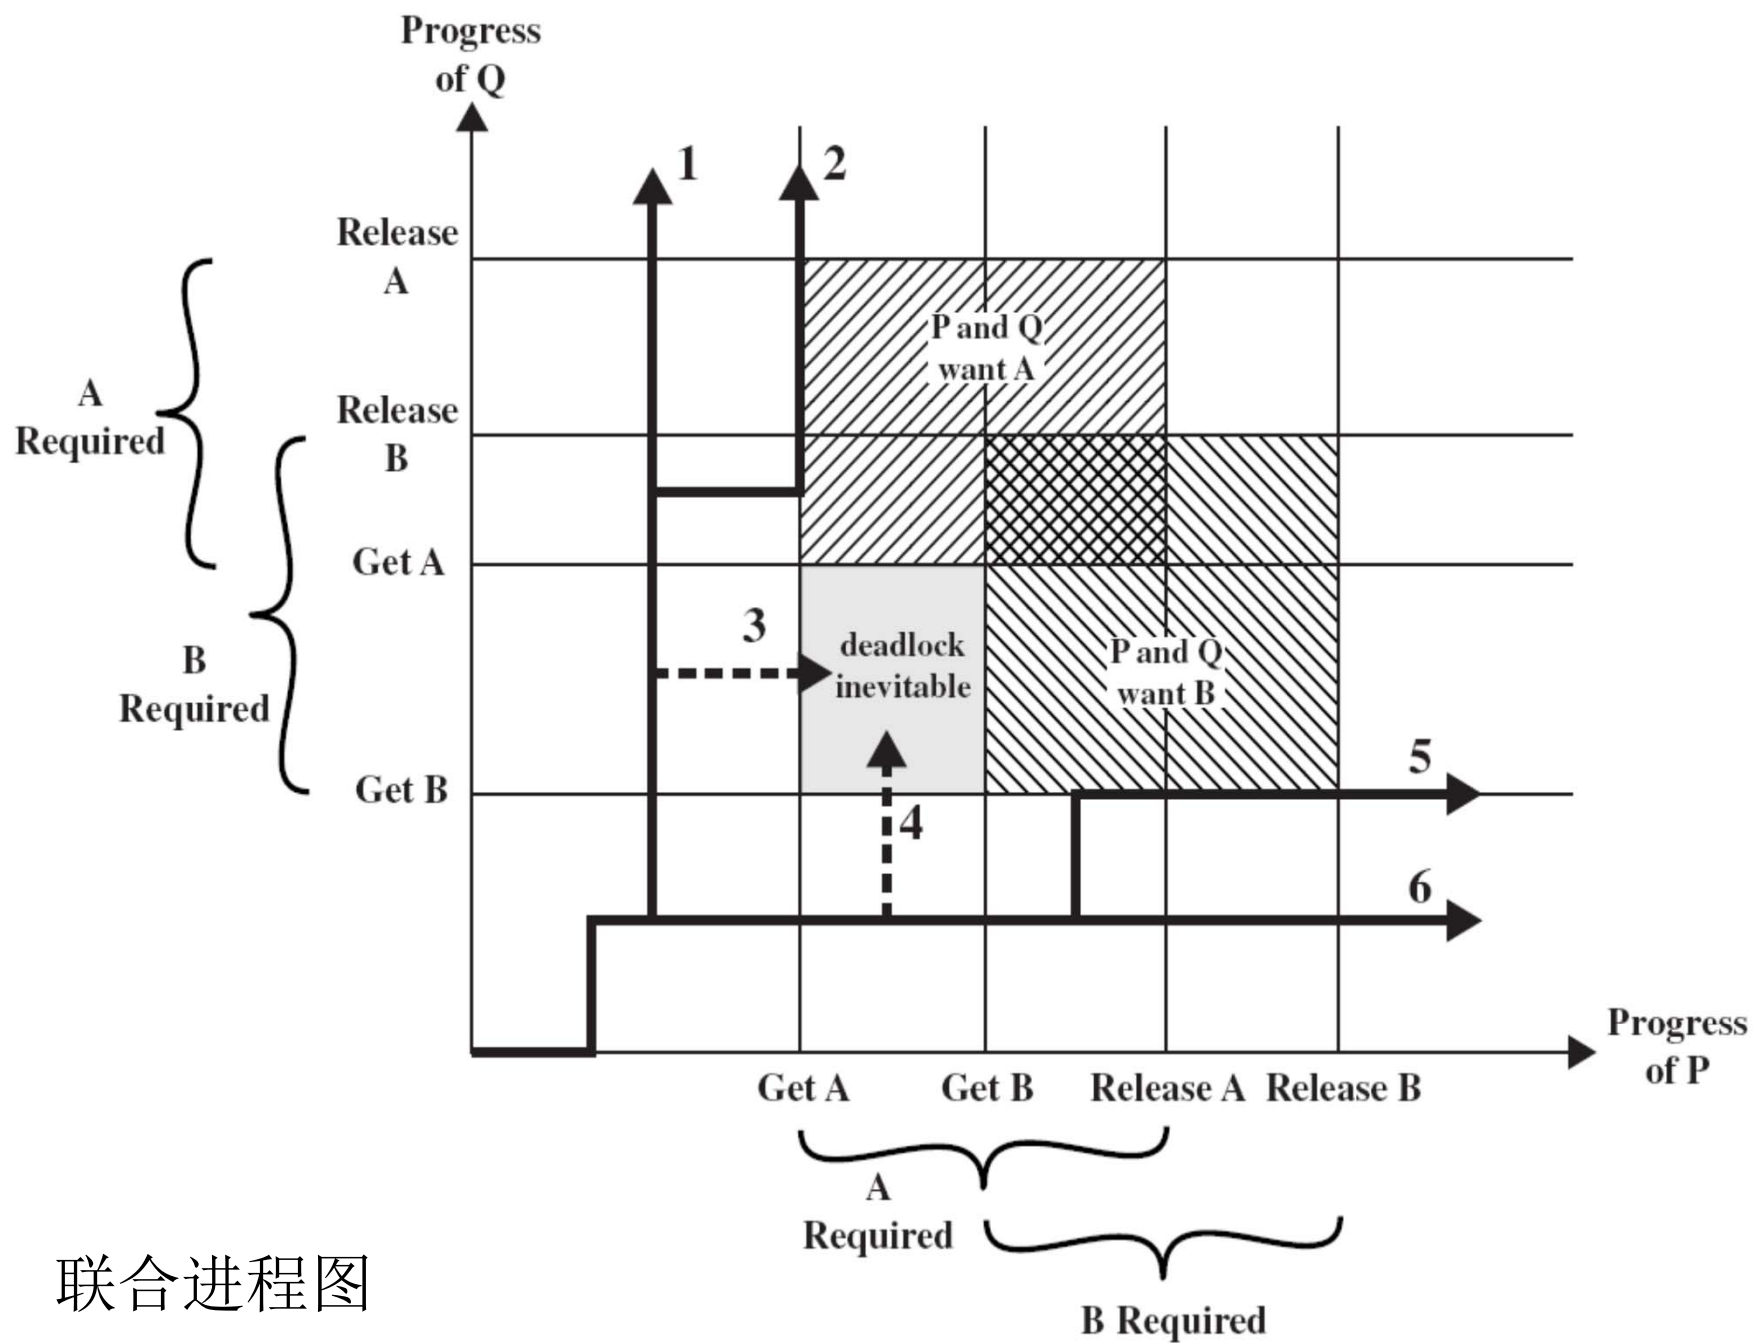
In 《modern operating systems》 3rd edition



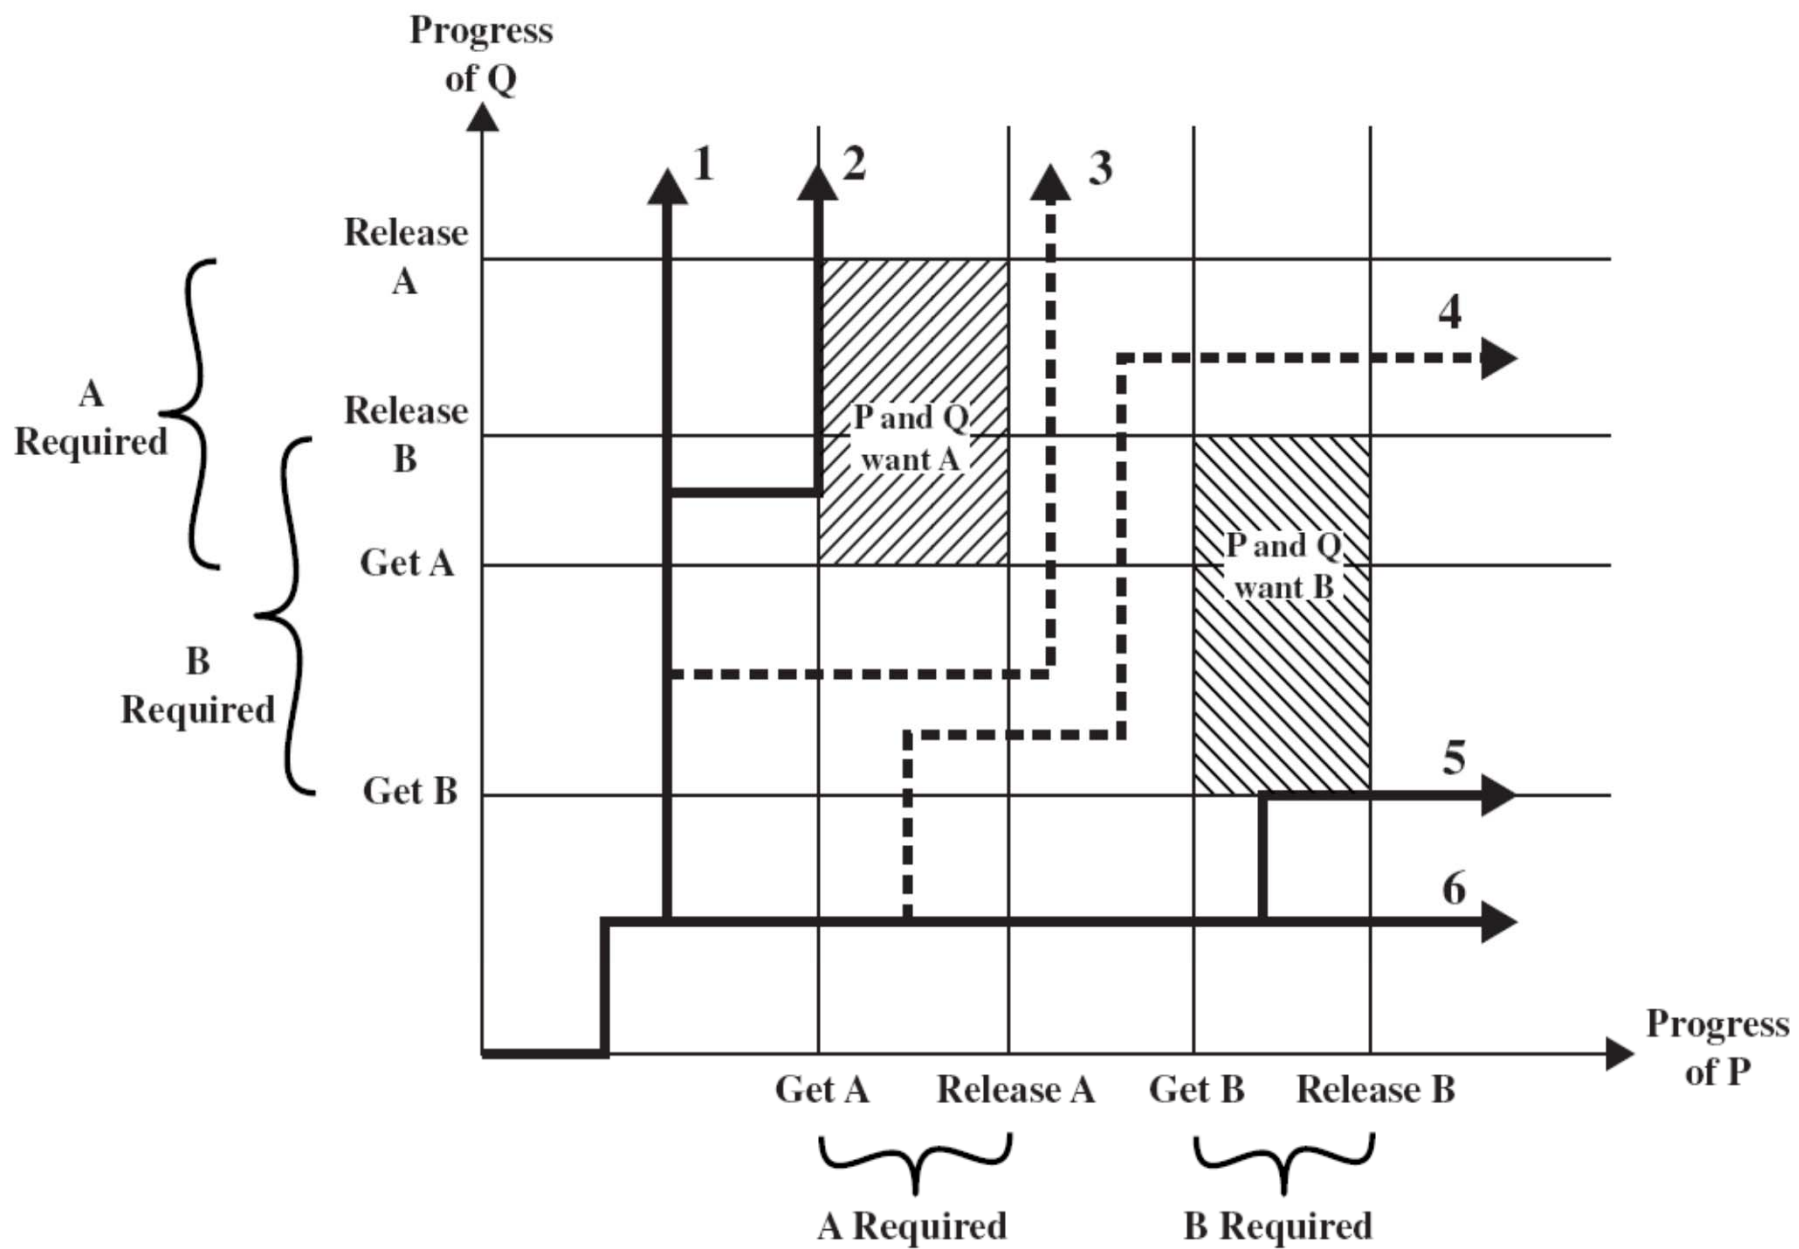
(a) Deadlock possible



(b) Deadlock



联合进程图



Two Categories of Resources

Reusable VS Consumable

- ▶ **Reusable Resources (可重用资源)**

- ▶ *A reusable resource is one that can be safely used by only one process at a time and is not depleted (销毁) by that use*

- ▶ **Consumable Resources (可消费资源)**

- ▶ *A consumable resource is one that can be created (produced) and destroyed (consumed)*



Reusable Resources (1)

- ▶ ***Processes obtain resources that they later release for reuse by other processes***
- ▶ ***Examples of reusable resources***
 - ▶ ***Processors, I/O channels, main and secondary memory, devices, data structures(files, databases, and semaphores)***



Reusable Resources (2)

Resource Acquisition 资源获得

- ▶ **Normally, a process may utilize a resource in only the following sequence:**

- ▶ **Request (or Get):** If the request cannot be granted immediately, then the requesting process must wait until it can acquire the resource.

- ▶ **Use:** The process can operate on the resource.

- ▶ **Release:** The process releases the resource

system calls: can be accomplished by semaphores



Reusable Resources (3)

Example of Deadlock

Process P

Request(D)
Request(T)
Perform function
Release(D)
Release(T)

Process Q



Request(T)
Request(D)
Perform function
Release(T)
Release(D)



Reusable Resources (4)

Another Example of Deadlock

- ▶ **Space is available for allocation of 200K bytes, and the following sequence of events occur**



- ▶ **Deadlock occurs if both processes progress to their second request**





Consumable Resources

- ▶ ***Created (produced) and destroyed (consumed) by a process***
- ▶ ***Examples of consumable resources***
 - ▶ ***Interrupts, signals, messages, and information in I/O buffers***



Consumable Resources Example of Deadlock

- ▶ Deadlock occurs if receive is blocking



Necessary Conditions for Deadlock (1)

- ▶ **Mutual Exclusion 互斥 (Resource)**
 - ▶ That is, only one process may use the resource at a time.
- ▶ **Hold-and-wait 占有且等待 (Process)**
 - ▶ A process may hold allocated resources while awaiting assignment of others




Necessary Conditions for Deadlock (2)

- ▶ **No preemption 非抢占 (Holding relationship)**
 - ▶ No resource can be forcibly removed from a process holding it. In other words, a resource can be released only voluntarily by the process holding it.
- ▶ **Circular wait 循环等待**
 - ▶ There is a closed chain of processes such that each process holds at least one resource needed by the next process in the chain.



Necessary Conditions for Deadlock (3)



These four conditions are not completely independent. For example, the circular-wait condition implies the hold-and-wait condition



Resource Allocation Graph (1)


资源分配图

- ▶ **A resource allocation graph consists of**
 - ▶ **A set of vertices V :**
 - ▶ The set of vertices V is partitioned into two subsets $P=\{P_1, P_2, \dots, P_n\}$ and $R=\{R_1, R_2, \dots, R_m\}$
 - ▶ **A set of edges E .**
 - ▶ A directed edge from process P_i to resource type R_j is called a **request edge** ($P_i \rightarrow R_j$)
 - ▶ A directed edge from resource type R_j to process P_i is called an **assignment edge** ($R_j \rightarrow P_i$)



Resource Allocation Graph (2)


- ▶ Each process P_i is represented as a circle
- ▶ Each resource type R_j as a square
 - ▶ Resource type R_j may have multiple instances. Each instance is represented as a dot within the square




Resource Allocation Graph (3)

- ▶ ***If the graph contains no cycles, then no process in the system is deadlocked***
- ▶ ***If the graph contains a cycle, then a deadlock may exist***
 - ▶ ***If each resource type has exactly one instance, then a cycle implies that a deadlock has occurred***
 - ▶ ***If each resource type has several instances, then a cycle does not necessarily imply that a deadlock has occurred***





Resource allocation graph with a deadlock



Resource allocation graph with a cycle but no deadlock

Four Strategies for Dealing with Deadlocks

- ▶ The Ostrich algorithm: Just ignore the problem
→ **Adopted by most operating systems**
- ▶ Detection and recovery: Let deadlock occur, detect them, and take action
- ▶ Dynamic avoidance: by carefully resource allocation
- ▶ Prevention: by structurally negating one of the four required conditions



Deadlock Prevention

Deadlock Prevention

- ▶ This strategy is to design a system such that the possibility of deadlock is excluded
- ▶ Prevention algorithms prevent deadlocks by *restraining how requests can be made*
 - ▶ The restraints ensure that at least one of the four necessary conditions can not hold



Deadlock Prevention Techniques (1)

- ▶ **Attacking “Mutual Exclusion”**
 - ▶ In general, mutual exclusion condition must hold for non-sharable resources
 - ▶ Mutual exclusion must be supported by OS if access to a resource requires mutual exclusion
 - ▶ In general, we can not prevent deadlocks by denying the mutual exclusion condition: Some resources are intrinsically nonsharable
- ▶ **Read-only files are a good example of a sharable resources**
 - ▶ A process never needs to wait for a sharable resource



Deadlock Prevention Techniques (2)

▶ **Attacking “Hold and Wait”**

- ▶ By requiring that a process *request all of its requested resources at one time* and blocking the process until all requests can be granted simultaneously
- ▶ Inefficient in two ways:
 - ▶ A process may wait a long time for all of its resources
 - ▶ Resources allocated may remain unused for a long period (but other processes can not use it)
- ▶ A process may not know in advance all of the resources required



Deadlock Prevention Techniques (3)

- ▶ **Attacking “No Preemption”**
 - ▶ If a process holding certain resources is denied a further request, that process must release its original resources, and if necessary, request them again together with the additional resource (主动放弃)
 - ▶ If a process requests a resource held by another process, the OS may preempt the second process and request it to release its resources (被抢占)
- ▶ **Limitation:** This approach is practical only for resources *whose state can be easily saved and restored later* (e.g. processors)



Deadlock Prevention Techniques (4)

- ▶ **Attacking “Circular Wait”**
 - ▶ Defining a linear ordering of all resource types
 - ▶ If a process has been allocated resource of type R , then it may request only those resources of types following R in the ordering
- ▶ **May be inefficient**



Deadlock Avoidance

Deadlock Avoidance

- ▶ A decision is made **dynamically** whether the current resource allocation request will, if granted, potentially lead to a deadlock
 - ▶ Thus, deadlock avoidance requires knowledge of future process resource requests



Two Approaches to Deadlock Avoidance

- ▶ **Process Initiation Denial** (进程启动拒绝)
 - ▶ *Do not start a process if its demands might lead to deadlock*
- ▶ **Resource Allocation Denial** (资源分配拒绝)
 - ▶ *Do not grant an incremental resource request to a process if this allocation might lead to deadlock*



Process Initiation Denial Problem Description

Resource = (R_1, R_2, \dots, R_m) \longrightarrow total amount of each resource

Available = (V_1, V_2, \dots, V_m) \longrightarrow total amount of each resource not allocated to a process

$$Claim = \begin{pmatrix} C_{11} & C_{12} & \Lambda & C_{1m} \\ C_{21} & C_{22} & \Lambda & C_{2m} \\ M & M & M & M \\ C_{n1} & C_{n2} & \Lambda & C_{nm} \end{pmatrix}$$

requirement of each process for each resource

$$Allocation = \begin{pmatrix} A_{11} & A_{12} & \Lambda & A_{1m} \\ A_{21} & A_{22} & \Lambda & A_{2m} \\ M & M & M & M \\ A_{n1} & A_{n2} & \Lambda & A_{nm} \end{pmatrix}$$

current allocation

Process Initiation Denial

Relationship Among These Matrices and Vectors

$$R_i = V_i + \sum_{k=1}^n A_{ki} \quad \text{for all } i: \text{ All resources are either available or allocated}$$

$$C_{ki} \leq R_i \quad \text{for all } k, i: \text{ No process can claim more than the total amount of resources in the system}$$

$$A_{ki} \leq C_{ki} \quad \text{for all } k, i: \text{ No process is allocated more resources of any type than the process originally claimed to need.}$$



Process Initiation Denial Policy

- ▶ **Refuse to start a new process if its resource requirements might lead to deadlock.**
- ▶ **A new process P_{n+1} can be started only if:**

$$R_i \geq C_{(n+1)i} + \sum_{k=1}^n C_{ki} \quad \text{for all } i$$

*It assumes the worst: all processes will make their maximum claims together
This strategy is hardly optimal!!*



Resource Allocation Denial

Description

- ▶ This strategy is also referred to as the banker's algorithm (银行家算法)
- ▶ State of the system is simply the current allocation of resources to process
 - ▶ The state consists of two vectors (Resource and Available), and two matrices (Claim and Allocation)
- ▶ **Safe or Unsafe?**
 - ▶ *Safe state* is where there is at least one sequence that does not result in deadlock
 - ▶ *Unsafe state* is a state that is not safe



(1)

Initial State

Which process can be run to completion with the resources available?

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim Matrix

	R1	R2	R3
P1	1	0	0
P2	6	1	2
P3	2	1	1
P4	0	0	2

Allocation Matrix

R1	R2	R3
9	3	6

Resource Vector

R1	R2	R3
0	1	1

Available Vector

(a) Initial state

P2!!

(2)

P2 Runs to Completion

	R1	R2	R3
P1	3	2	2
P2	0	0	0
P3	3	1	4
P4	4	2	2

Claim Matrix

	R1	R2	R3
P1	1	0	0
P2	0	0	0
P3	2	1	1
P4	0	0	2

Allocation Matrix

R1	R2	R3
6	2	3

Available Vector

(b) P2 runs to completion



(3)

P1 Runs to Completion

After choosing P1 to be run to completion, the state is as follows:

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	3	1	4
P4	4	2	2

Claim Matrix

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	2	1	1
P4	0	0	2

Allocation Matrix

R1	R2	R3
7	2	3

Available Vector

(c) P1 runs to completion



(4)

P3 Runs to Completion

Next, we can complete P3

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	4	2	2

Claim Matrix

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	0	0	2

Allocation Matrix

R1	R2	R3
9	3	4

Available Vector

(d) P3 runs to completion

Finally, we can complete P4. All the 4 processes can be run to completion. Thus, the initial state is safe.



Resource Allocation Denial Solution

- ▶ **One possible deadlock avoidance strategy is to ensure that the system of processes and resources is *always* in a safe state**
- ▶ **When a process makes a request for a set of resources,**
 - ▶ Assume that the request is granted
 - ▶ Update the system state accordingly
 - ▶ And then determine whether the result is a safe state
 - ▶ If safe, grant the request; otherwise, block the process until it is safe to grant the request.



Determination of an Unsafe State

Consider the following safe initial state

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim Matrix

	R1	R2	R3
P1	1	0	0
P2	5	1	1
P3	2	1	1
P4	0	0	2

Allocation Matrix

R1	R2	R3
9	3	6

Resource Vector

R1	R2	R3
1	1	2

Available Vector

Suppose that P2 makes a request for one additional unit of R1 and one additional unit of R3. Should this request be granted?

► *How about P1 making the same request instead?*

Determination of an Unsafe State

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim Matrix

	R1	R2	R3
P1	2	0	1
P2	5	1	1
P3	2	1	1
P4	0	0	2

Allocation Matrix

R1	R2	R3
0	1	1

Available Vector

(b) P1 requests one unit each of R1 and R3



Deadlock Avoidance

Advantages and Restrictions

▶ **Advantages**

- ▶ It is not necessary to preempt or rollback processes (as in deadlock detection)
- ▶ It is less restrictive than deadlock prevention

▶ **Restrictions on its use:**

- ▶ Maximum resource requirement must be stated in advance
- ▶ Processes under consideration must be independent; no synchronization requirements
- ▶ There must be a fixed number of resources to allocate
- ▶ No process may exit while holding resources



Deadlock Detection and Recovery

Deadlock Detection and Recovery

- ▶ If a system does not employ either a deadlock-prevention or a deadlock-avoidance algorithm, then a deadlock may occur.
- ▶ System must provide:
 - ▶ An algorithm that *determines whether a deadlock has occurred*
 - ▶ An algorithm to *recover from the deadlock*



Problem Description

- ▶ **Allocation matrix and Available vector**
 - ▶ As described in deadlock avoidance
- ▶ **Request matrix Q**
 - ▶ q_{ij} represents the amount of resources of type j requested by process i .

Note: The Claim matrix is not needed in deadlock detection, while it is a must in deadlock avoidance algorithms.



Deadlock Detection Algorithm

- ▶ **Step 1:** Mark each process that has a row in the Allocation Matrix of all zeros
- ▶ **Step 2:** Initialize a temporary vector W to equal the Available vector
- ▶ **Step 3:** Find an index i such that process i is currently unmarked and the i th row of Q is less than or equal to W . If no such row is found, terminate the algorithm
- ▶ **Step 4:** If such a row is found, mark process i and add the corresponding row of the allocation matrix to W . Return to Step 3.



Example for Deadlock Detection

	R1	R2	R3	R4	R5
P1	0	1	0	0	1
P2	0	0	1	0	1
P3	0	0	0	0	1
P4	1	0	1	0	1

Request Matrix Q

	R1	R2	R3	R4	R5
P1	1	0	1	1	0
P2	1	1	0	0	0
P3	0	0	0	1	0
P4	0	0	0	0	0

Allocation Matrix A

R1	R2	R3	R4	R5
2	1	1	2	1

Resource Vector

R1	R2	R3	R4	R5
0	0	0	0	1

Available Vector

P3 is less or equal to W
P4 has the row of all zero

W	0	0	0	0	1
---	---	---	---	---	---

Initialize W to equal the Available Vector

W	0	0	0	1	1
---	---	---	---	---	---

Add the row of P3 in Allocation matrix to W

No other unmarked process has a row in Q that is less than or equal to W!!!!
Deadlock has occurred!!

When to Detect Deadlocks?

- ▶ ***There are several candidate strategies:***
 - ▶ ***To check every time a resource request is made (certain to detect deadlocks as early as possible, but potentially expensive in terms of CPU time)***
 - ▶ ***To check every k minutes***
 - ▶ ***To check only when the CPU utilization has dropped below some threshold (if enough processes are deadlocked, there will be few runnable processes)***



Recovery from Deadlock

- ▶ **How to deal with it once a deadlock has been detected?**
 - ▶ To inform the operator that a deadlock has occurred, and let the operator to deal with it **manually**; or
 - ▶ To let the system recover from the deadlock **automatically**.
 - ▶ **Process termination**: to abort one or more processes to break the circular wait
 - ▶ **Resource preemption**: to preempt some resources from one or more of the deadlocked process



Recovery Strategies once Deadlock Detected

- ▶ **Abort all deadlocked processes**
 - ▶ The simplest and the most common
- ▶ **Back up each deadlocked process to some previously defined checkpoint (检查点), and restart all process**
 - ▶ Rollback (回滚) and Restart (重启) mechanisms should be built into system
 - ▶ original deadlock may occur
- ▶ **Successively abort deadlocked processes until deadlock no longer exists**
- ▶ **Successively preempt resources until deadlock no longer exists**
 - ▶ A process that has a resource preempted from it must be rolled back to a point prior to its acquisition of that resource.



Selection Criteria

- ▶ **Which process to abort or resources from which process to be preempted?**
 - ▶ Least amount of processor time consumed so far
 - ▶ Least amount of output produced so far
 - ▶ Most estimated time remaining
 - ▶ Least total resources allocated so far
 - ▶ Lowest priority
 - ▶ Interactive or batch




The Ostrich Algorithm

The Ostrich Algorithm

- ▶ ***The basic idea: stick your head in the sand and pretend there is no problem at all.***
 - ▶ *It seems not to be a viable approach, but it is used in some operating systems (including UNIX)*
- ▶ ***How often does a deadlock occur? How serious is a deadlock?***




Safe State, Unsafe State, and Deadlocked State



Dining Philosophers Problem

Dining Philosophers Problem



A First Solution to the Dining Philosophers Problem

```
semaphore fork[5]={1};
int i;
void philosopher(int i){
    while(true){
        think();
        wait(fork[i]);
        wait(fork[(i+1) mod 5]);
        eat();
        signal(fork[(i+1) mod 5]);
        signal(fork[i]);
    }
}
void main(){
    parbegin(philosopher(0), philosopher(1), philosopher(2),
             philosopher(3), philosopher(4));
}
```

**This is a failed
solution!!**

**How to use deadlock
prevention to solve it?**

A Second Solution to the Dining Philosophers Problem

```
semaphore fork[5]={1};
semaphore room={4};
int i;
void philosopher(int i){
    while(true){
        think();
        wait(room);
        wait(fork[i]);
        wait(fork[(i+1) mod 5]);
        eat();
        signal(fork[(i+1) mod 5]);
        signal(fork[i]);
        signal(room);
    }
}
```

At most 4 philosophers are allowed to get into the dining room at a time

At least one philosopher will have access to two forks.

