




编译原理/Compiler Instruction Selection


周雅倩
复旦大学计算机科学技术学院
zhouyagqian@fudan.edu.cn
2018/12/13



References


- <http://www.stanford.edu/class/cs143/>
 - Partial contents are copied from the slides of Alex Aiken

2



Backend Optimizations

- Instruction selection**
 - Translate low-level IR to assembly instructions
 - A machine instruction may model multiple IR instructions
 - Especially applicable to CISC architectures
- Register Allocation(chapter 11)**
 - Place variables into registers
 - Avoid spilling variables on stack




Outline

- [Tree Patterns](#)
- [Algorithms For Instruction Selection](#)
- [Low-level IR to IR Tree/DAG](#)



TREE PATTERNS



Tree patterns


- Each IR tree node expresses one operation.
 - Most real architectures can perform several in one instruction as at right.
 - Instruction selection - determining machine instructions to implement an IR tree
 - Tree pattern - the IR tree fragment that corresponds to a valid instruction
- Tiling - instruction selection to cover IR tree with minimal tree patterns

| Name | Effect | Trees |
|-------|---------------------------------|-------|
| ADD | $r_i \leftarrow r_j + r_k$ | TEMP |
| MUL | $r_i \leftarrow r_j \times r_k$ | |
| SUB | $r_i \leftarrow r_j - r_k$ | |
| DIV | $r_i \leftarrow r_j / r_k$ | |
| ADDI | $r_i \leftarrow r_j + c$ | |
| SUBI | $r_i \leftarrow r_j - c$ | |
| LOAD | $r_i \leftarrow M[r_j + c]$ | |
| STORE | $M[r_j + c] \leftarrow r_i$ | |
| MOVEM | $M[r_j] \leftarrow M[r_i]$ | |

FIGURE 9.1. Arithmetic and memory instructions. The notation $M[x]$ denotes the memory word at address x .

Example

- Find a tiling that cover the IR tree at right.
- Give the effect of the tilings.
- Are other tilings possible?



1. addi r1 ← r0 + 1
2. load r2 ← M[r1 + 2]

1. addi r1 ← r0 + 1
2. addi r2 ← r0 + 2
3. add r3 ← r1 + r2
4. load r4 ← M[r3 + 0]

Example

- $a[i] = x$
 - where i is register, a and x are frame-resident
 - a is frame offset to reference to an array




FIGURE 9.2. A tree tiled in two ways.

One tile covers one node

ADDI $r_1 \leftarrow r_0 + a$
ADD $r_1 \leftarrow \text{fp} + r_1$
LOAD $r_1 \leftarrow M[r_1 + 0]$
ADDI $r_2 \leftarrow r_0 + 4$
MUL $r_2 \leftarrow r_1 \times r_2$
ADD $r_1 \leftarrow r_1 + r_2$
ADDI $r_2 \leftarrow r_0 + x$
ADD $r_2 \leftarrow \text{fp} + r_2$
LOAD $r_2 \leftarrow M[r_2 + 0]$
STORE $M[r_1 + 0] \leftarrow r_2$

Optimal and Optimum Tilings

- least cost - fewest instructions, least time, etc.
- Optimum (最优) - tiling whose cost sum to lowest possible value
- Optimal (最佳) - no two adjacent tiles can be combined into a single tile of lower total cost.

Optimal is not necessarily optimum

- Suppose all instructions have cost of 1 except MOVEM which has cost of m .
- $m > 1$ then (a) optimum, (b) optimal:
 - (a) = 6
 - (b) > 6
- $m < 1$ then (a) optimal, (b) optimum:
 - (a) = 6
 - (b) < 6
- $m = 1$ then (a) and (b) are optimum:
 - (a) = (b)

ALGORITHMS FOR INSTRUCTION SELECTION

Algorithms For Instruction Selection

- CISC - Complex Instruction Set
Computers often accomplish multiple operations per instruction, can have large tilings
- RISC - Reduced Instruction Set
Computers generally accomplish few operations per instruction, have smaller tilings
- Optimum/Optimal differences larger on CISC than on RISC

Maximal Munch - Optimal tiling

- Algorithm
 - Start at root
 - Find largest tile that fits (covers the most nodes) leaving subtrees
 - Generate instruction for tiling
 - Repeat on each subtree


Implementation

- Two recursive functions
 - munchStm for statements
 - One clause for each statement
 - munchExp for expressions
 - each clause will match one tile
 - clauses ordered by tile preference (biggest to smallest)
- Key is finding target instruction that covers the largest tiling
- Order into cases where largest tiling matches are checked first.
- Jouette STORE and MOVEM same size tiling, choose either
- Top-down, depth-first recursion over IR tree, emit instructions *after* visiting subtrees
- Only instruction is emitted, registers allocated later

Exercise 2

- Munch the following IR into Jouette instructions using PROGRAM 9.3

```
MEM(BINOP(PLUS, e1, CONST(2)))  
MOVE(TEMP(t0), CONST(3))
```



Dynamic Programming

- Assigns a **cost** to every node in tree, the sum of the cost of subtree nodes
- Bottom-up, summing cost of node n subtrees first
- Matches n with each tile kind, minimum cost match selected
- $\text{cost}(n) = \text{sum}(\text{subtree cost}) + \text{cost}(\text{tiling } n)$
- leaves** - places where subtrees can be attached
- Each tile has 0 or more leaves

Example

| Tile | Instruction | Tile Cost | Leaves Cost | Total Cost |
|------|-------------|-----------|-------------|------------|
| | ADD | 1 | 1+1 | 3 |
| | ADDI | 1 | 1 | 2 |
| | ADDI | 1 | 1 | 2 |

Example - Cost is number of tilings

- $\text{cost}(\text{ADD}) = 3$ since total 3 the number of tilings required to cover ADD and subtrees
 - $\text{ADDI } r_2 \leftarrow r_0 + 2$
 - $\text{ADDI } r_1 \leftarrow r_0 + 1$
 - $\text{ADD } r_3 \leftarrow r_1 + r_2$
- $\text{cost}(\text{ADDI}) = 2$ since total 2 the number of tilings required to cover ADDI and subtrees
 - $\text{ADDI } r_2 \leftarrow r_0 + 2$
 - $\text{ADDI } r_3 \leftarrow r_2 + 1$

Using minimum cost of subtree

| Tile | Instruction | Tile Cost | Leaves Cost | Total Cost |
|------|-------------|-----------|-------------|------------|
| | LOAD | 1 | 2 | 3 |
| | LOAD | 1 | 1 | 2 |
| | LOAD | 1 | 1 | 2 |

$\text{ADDI } r_1 \leftarrow r_0 + 1$
 $\text{LOAD } r_1 \leftarrow M[r_1 + 2]$

- $\text{cost}(\text{MEM}) = 3$
 - $\text{ADDI } r_2 \leftarrow r_0 + 2$
 - $\text{ADDI } r_3 \leftarrow r_2 + 1$
 - $\text{LOAD } r_4 \leftarrow M[r_3 + 0]$
- $\text{cost}(\text{MEM}) = 2$
 - $\text{ADDI } r_2 \leftarrow r_0 + 2$
 - $\text{LOAD } r_4 \leftarrow M[r_2 + 1]$
- $\text{cost}(\text{MEM}) = 2$
 - $\text{ADDI } r_1 \leftarrow r_0 + 1$
 - $\text{LOAD } r_4 \leftarrow M[r_1 + 2]$

•Note that other costs are possible:
 •Number of memory accesses
 •Number of registers used

Tree Grammars


- Use a context-free grammar to describe the tiles
 - nonterminals s (statements), a (expressions)
- Skip for now

Fast Matching

- Maximal munch and dynamic-programming algorithms examine all tiles that match each node.
- Tile match if each non-leaf node of tile is labeled with same operator as node of tree (MEM, CONST, etc.)
- More efficient to examine only tiles with matching root label

Example

- The MEM root has only 4 possible subtrees, examine only those rather than recursively examining every possible expression again.
- In practice, *visitor* pattern yields similar performance; discussed at end of notes.



Efficiency of Tiling Algorithms

- Linear runtime of both Maximal Munch and Dynamic Programming
- Dependent on number of nodes in input tree
- In practice
 - Instruction selection is very fast
 - Even lexical analysis takes more time to run

MM vs DP

- Maximal Munch
 - Good for RISC
 - Tile size is small and of uniform cost
 - Usually no difference between optimum and optimal tiling
- Dynamic Programming
 - Good for CISC
 - Tile size is large and of widely varying cost
 - Small, but noticeable differences between optimum and optimal tiling

CISC&RISC

CISC&RISC

| CISC | RISC |
|---|---|
| Emphasis on hardware | Emphasis on software |
| Includes multi-clock complex instructions | Single-clock, reduced instruction only |
| Memory-to-memory: "LOAD" and "STORE" incorporated in instructions | Register to register: "LOAD" and "STORE" are independent instructions |
| Small code sizes, high cycles per second | Low cycles per second, large code sizes |
| Transistors used for storing complex instructions | Spends more transistors on memory registers |

Complex Instruction Set Computers

Reduced Instruction Set Computers

Example

x = x + 1

load r1 ← M[fp+4]
addi r2 ← r1 + 1
store M[fp+4] ← r2

CISC
add[fp+4], \$1

IR Tree/DAG

```

graph TD
    MOVE --> MEM1[MEM]
    MOVE --> PLUS1[+]
    MOVE --> CONST1[Const 1]
    MEM1 --> PLUS2[+]
    PLUS2 --> FP1[FP]
    PLUS2 --> CONST4_1[Const 4]
    PLUS1 --> MEM2[MEM]
    PLUS1 --> PLUS3[+]
    PLUS3 --> FP2[FP]
    PLUS3 --> CONST4_2[Const 4]
  
```

LOW-LEVEL IR TO IR TREE/DAG

Instruction Selection

- Different sets of instructions in low-level IR and in the target machine
- Instruction selection** = translate low-level IR to assembly instructions on the target machine
- Straightforward solution**: translate each low-level IR instruction to a sequence of machine instructions
- Example:

$x = y + z$

→

```

mov y, r1
mov z, r2
add r2, r1
mov r1, x

```

Instruction Selection

- Problem**: straightforward translation is inefficient
 - One machine instruction may perform the computation in multiple low-level IR instructions
 - Excessive memory traffic

Example

- Consider a machine that includes the following **instructions**:

| | |
|---|--|
| <pre> add r2, r1 mulc c, r1 load r2, r1 store r2, r1 movem r2, r1 movex r3, r2, r1 </pre> | <pre> r1 ← r1 + r2 r1 ← r1 * c r1 ← *r2 *r1 ← r2 *r1 ← *r2 *r1 ← *(r2 + r3) </pre> |
|---|--|

Example

- Consider the computation: $a[i+1] = b[j]$
- Assume **a, b, i, j** are global variables
 - register **ra** holds address of **a**
 - register **rb** holds address of **b**
 - register **ri** holds value of **i**
 - register **rj** holds value of **j**

Low-level IR:

```

t1 = j * 4
t2 = b + t1
t3 = *t2
t4 = i + 1
t5 = t4 * 4
t6 = a + t5
*t6 = t3

```


Possible Translation

- Address of b[j]: `mulc 4, rj`
`add rj, rb`
- Load value b[j]: `load rb, r1`
- Address of a[i+1]: `add 1, ri`
`mulc 4, ri`
`add ri, ra`
- Store into a[i+1]: `store r1, ra`

Low-level IR:

```

t1 = j*4
t2 = b+t1
t3 = *t2
t4 = i+1
t5 = t4*4
t6 = a+t5
*t6 = t3
  
```

Another Translation

- Address of b[j]: `mulc 4, rj`
`add rj, rb`
- Address of a[i+1]: `add 1, ri`
`mulc 4, ri`
`add ri, ra`
- Store into a[i+1]: `movem rb, ra`

Low-level IR:

```

t1 = j*4
t2 = b+t1
t3 = *t2
t4 = i+1
t5 = t4*4
t6 = a+t5
*t6 = t3
  
```

Yet Another Translation

- Index of b[j]: `mulc 4, rj`
- Address of a[i+1]: `add 1, ri`
`mulc 4, ri`
`add ri, ra`
- Store into a[i+1]: `movex rj, rb, ra`

Low-level IR:

```

t1 = j*4
t2 = b+t1
t3 = *t2
t4 = i+1
t5 = t4*4
t6 = a+t5
*t6 = t3
  
```

Issue: Instruction Costs

- Different machine instructions have different costs
 - Time cost: how fast instructions are executed
 - Space cost: how much space instructions take
- Example: **cost = number of cycles**
 - `add r2, r1` cost=1
 - `mulc c, r1` cost=10
 - `load r2, r1` cost=3
 - `store r2, r1` cost=3
 - `movem r2, r1` cost=4
 - `movex r3, r2, r1` cost=5
- Goal: find translation with smallest cost

How to Solve the Problem?

- Difficulty: low-level IR instruction matched by a machine instructions may not be adjacent
- Example: `movem rb, ra`
- Idea: use tree-like representation!
 - Easier to detect matching instructions

Low-level IR:

```

t1 = j*4
t2 = b+t1
t3 = *t2
t4 = i+1
t5 = t4*4
t6 = a+t5
*t6 = t3
  
```

Tree Representation

- Goal: determine parts of the tree that correspond to machine instructions

Low-level IR:

```

t1 = j*4
t2 = b+t1
t3 = *t2
t4 = i+1
t5 = t4*4
t6 = a+t5
*t6 = t3
  
```

Tree diagram for `a[i+1] = b[j]`:

```

graph TD
    store --> plus1[+]
    store --> load[load]
    plus1 --> a[a]
    plus1 --> mulc[*]
    mulc --> plus2[+]
    plus2 --> i[i]
    plus2 --> 1[1]
    mulc --> 4[4]
    load --> plus3[+]
    plus3 --> b[b]
    plus3 --> mul4[*]
    mul4 --> j[j]
    mul4 --> 4[4]
  
```

Tiles

■ Tile = tree patterns (subtrees) corresponding to machine instructions

Low-level IR:

```

t1 = j*4
t2 = b+t1
t3 = *t2
t4 = i+1
t5 = t4*4
t6 = a+t5
*t6 = t3
  
```

movem rb, ra

Tiling

■ Tiling = cover the tree with disjoint tiles

Low-level IR:

```

t1 = j*4
t2 = b+t1
t3 = *t2
t4 = i+1
t5 = t4*4
t6 = a+t5
*t6 = t3
  
```

movem rb, ra

Tiling

store rb, ra

movem rj, rb, ra

Directed Acyclic Graphs

■ Tree representation: appropriate for instruction selection

- Tiles = subtrees → machine instructions

■ DAG = more general structure for representing instructions

- Common sub-expressions represented by the same node
- Tile the expression DAG

Example


```

t = y+1
y = z*t
t = t+1
z = t*y
  
```

Big Picture


■ What the compiler has to do:

1. Translate low-level IR code into DAG representation
2. Then find a good tiling of the DAG
 - Maximal munch algorithm
 - Dynamic programming algorithm







DAG Construction

- **Input:** a sequence of low IR instructions in a basic block
- **Output:** an expression DAG for the block
- **Idea:**
 - Label each DAG node with variable holding that value
 - Build DAG bottom-up
 - A variable may have multiple values in a block
- Use different variable indices for different values of the variable: t_0, t_1, t_2 , etc



编译原理











Algorithm


```

index[v] = 0 for each variable v
For each instruction I (in the order they appear)
  For each variable v that I directly uses, with
    n=index[v]
    if node vn doesn't exist
      create node vn, with label vn
  Create expression node for instruction I, with
    children
      { vn | v ∈ use[I] }
  For each v ∈ def[I]
    index[v] = index[v] + 1
  If I is of the form x = ... and n = index[x]
    label the new node with xn
          
```




编译原理



Issues

- **Function/method calls**
 - May update global variables or object fields
 - $def[I]$ = set of globals/fields
- **Store instructions**
 - May update any variable
 - If stack addresses are not taken (e.g., Java), $def[I]$ = set of heap objects



编译原理

