

复旦大学

编译原理/Compiler Lexical Analysis

周雅倩
复旦大学计算机科学技术学院
{zhouyaqian}@fudan.edu.cn
2018/10/11

References

- Partial contents are copied from the slides of Alex Aiken
 - <http://www.stanford.edu/class/cs143/>
 - <https://www.coursera.org/course/compilers>
- http://en.wikipedia.org/wiki/Comparison_of_parser_generators

2

Outline

- Specifying lexical structure using regular expressions
- Finite automata
 - Deterministic Finite Automata (DFAs)
 - Non-deterministic Finite Automata (NFAs)
- Implementation of regular expressions
 - RegExp => NFA => DFA => Tables

Outline

- 语言 (Language)
- 正则表达式 (Regular Expressions)
- 有限自动机 (Finite Automata)

Definitions

- The **lexical analyzer** (词法分析器) produces a certain **token** (单词) wherever the input contains a string of characters in a certain set of strings.
- The set of strings is described by a **rule** (规则).
- A **pattern** (模式) is associated with a token.
- A **lexeme** (词素) is a sequence of characters matching the pattern.

Token	Sample lexemes	Informal description of pattern
const	const	const
if	if	if
relation	<,<=,>,>=	< or <= or > or >=
id	pi, count, D2	Letter followed by letters or digits
num	3.1416, 0, 6.02E23	Any numeric constant
literal	"core dumped"	Any characters between "and" except ""

Some Language aspects

- Fortran requires certain constructs in certain positions of input line, complicating lexical analysis.
- Modern languages – free-form input
 - Position in an input line is not important.

Some Language aspects

- Sometimes blanks are allowed within lexemes. For e.g. in Fortran X VEL and XVEL represent the same variable.

e.g. `DO 5 I = 1,25` is a Do- statement while
`DO 5 I = 1.25` is an assign statement.

Lexical Analysis in PL/I

- Most languages reserve keywords. Some languages do not, thus complicating lexical analysis.
 - e.g. PL/I

`IF THEN THEN THEN = ELSE; ELSE ELSE = THEN;`

- PL/I Declarations:
 - `DECLARE (ARG1, . . . , ARGN)`
- Can't tell whether `DECLARE` is a keyword or array reference until after the `)`.
 - Requires arbitrary look ahead!

Lexical Analysis in C++

- Unfortunately, the problems continue today
- C++ template syntax:
 - `Foo<Bar>`
- C++ stream syntax:
 - `cin >> var;`
- But there is a conflict with nested templates:
 - `Foo<Bar<Bazz>>`

Review

- The goal of **lexical analysis** is to
 - Partition the input string into lexemes
 - Identify the token of each lexeme
- Left-to-right scan => lookahead sometimes required

Attributes of Tokens

- If two or more lexemes match the pattern for a token then the lexical analyzer must provide additional information with the token.
- Additional information is placed in a symbol-table entry and the lexical analyzer passes a pointer/reference to this entry.
- E.g. The Fortran statement:

`E = M * C ** 2`

It has 7 tokens and associated attribute-values:

- `<id, reference to symbol-table entry for E>`
- `<assign_op, >`
- `<id, reference to symbol-table entry for M>`
- `<mult_op, >`
- `<id, reference to symbol-table entry for C>`
- `<exp_op, >`
- `<num, integer value 2>`

Sample of objects in Lex

Token
-index (type) : int
-stringVal : string(idl)
-numVal : float(idl)
-line : unsigned long(idl)
-charBegin : unsigned short(idl)
-charEnd

SymbolTable
+insert() : Token
+delete() : Token
+query() : Token

Errors in Lex

- Unmatched pattern
- Simplest -> Panic mode
- Recovery and continue on
- Many times, lex has only localized view
- E.g. fi (a == f(x))

Input Buffering

- Buffer pairs: The input buffer has two halves with N characters in each half.
- N might be the size of a disk block like 1024 or 4096.
- Mark the end of the input stream with a special character eof. Maintain two pointers into the buffer marking the beginning and end of the current lexeme.

E = M * | C * * 2 eof
 lexeme_start forward

Specifying Formal Languages

- Two Dual Notions
 - Generative approach (grammar or regular expression)
 - Recognition approach (automaton)
- Many theorems to transform one approach automatically to another

正则表达式 REGULAR EXPRESSIONS

Language

- Language denotes any set of strings over alphabets (very broad definition)

Examples of Languages

- **Alphabet** = English characters
- **Language** = English sentences
 - Not every string of English characters is an English sentence
- **Alphabet** = ASCII
- **Language** = C programs
 - Note: ASCII character set is different from English character set

Specifying Tokens

- String is a finite sequence of symbols
 - E.g. tech is a string length of four
 - The empty string, denoted ϵ , is a special string length of zero
 - Prefix of s
 - Suffix of s
 - Substring of s
 - proper prefix, suffix, substring of s ($x \neq s$, $x \neq \epsilon$) (真前缀, 后缀, 子串)
 - Subsequence of s

联结 (Concatenation)

- If x and y are strings then the **concatenation** of x and y , written xy , is the string formed by appending y to x .
 - If $x = \text{dog}$ and $y = \text{house}$ then $xy = \text{doghouse}$.
- **String Exponentiation**: If x is a string then $x^2 = xx$, $x^3 = xxx$, etc. $x^0 = \epsilon$.
 - If $x = \text{ba}$ and $y = \text{na}$ then $xy^2 = \text{banana}$.

Specifying Tokens

- Abstract languages like the empty set $\{\epsilon\}$, the set only empty strings
- Languages Operations
 - Union (并)
 - Concatenation (联结)
 - Closure (闭包)

Language Operations

- **Union**
 - $L \cup M = \{s \mid s \text{ is in } L \text{ or in } M\}$
 - If L and M are languages then $L \cup M$ is the language containing all strings in L and all strings in M .
- **Concatenation**
 - $\{st \mid s \text{ is in } L \text{ and } t \text{ is in } M\}$
 - If L and M are languages then LM is the language that contains concatenations of any string in L with any string in M .
- **Kleene Closure**
 - If L is a language then $L^+ = \{\epsilon\} \cup L \cup LL \cup LLL \cup LLLL \cup \dots$
- **Positive Closure (正闭包)**
 - If L is a language then $L^+ = L \cup LL \cup LLL \cup LLLL \cup \dots$

Language Operations

- For e.g. Let $L = \{A, B, \dots, Z, a, b, \dots, z\}$ and let $D = \{0, 1, 2, \dots, 9\}$
 - $L \cup D$?
 - LD ?
- $L^4 = LLLL$ is the set of all four-letter strings,
- L^+ is the set of all strings of letters including the empty string, ϵ ,
- $L(L \cup D)^*$ is the set of all strings of letters and digits that begin with a letter, and
- D^+ is the set of all strings of one or more digits.

Regular Expressions

- **Regular expressions (正则表达式)** are an important notation for specifying patterns.
 - Letter (letter| digit)* → ?
- **Alphabet (字母表)**: A finite set of symbols.
 - {0,1} is the binary alphabet.
 - ASCII and EBCDIC are two examples of computer alphabets.
- A **string (字符串)** over an alphabet is a finite sequence of symbols drawn from that alphabet.
 - 011011 is a string of length 6 over the binary alphabet.
 - The empty string denoted ϵ , is a special string of length zero.
- A **language (语言)** is any **set of strings** over some fixed alphabet.

Rules for Regular Expressions Over Alphabet Σ

- ϵ is a regular expression denoting $\{\epsilon\}$, the set containing the empty string.
- If a is a symbol in Σ then a is a regular expression denoting $\{a\}$.
- If r and s are regular expressions denoting languages $L(r)$ and $L(s)$, respectively then:
 - $r|s$ is a regular expression denoting $L(r) \cup L(s)$,
 - rs is a regular expression denoting $L(r)L(s)$,
 - r^* is a regular expression denoting $(L(r))^*$.

Precedence

- The unary operator $*$ has highest precedence
- **Concatenation** has second highest precedence
- $|$ has the lowest precedence.
- All operators are **left associative**.

E.g. Pascal Identifiers

letter $\rightarrow A|B|...|Z|a|b|...|z$
 digit $\rightarrow 0|1|...|9$
 id $\rightarrow \text{letter} (\text{letter} | \text{digit})^*$

E.g. Unsigned Numbers in Pascal

digits $\rightarrow \text{digit} \text{ digit}^*$
 opt_frac $\rightarrow \text{.digits} | \epsilon$
 opt_exp $\rightarrow (E(+ | - | \epsilon) \text{ digits}) | \epsilon$
 num $\rightarrow \text{digits opt_frac opt_exp}$

Shorthand Notations

- If r is a regular expression then :
 - r^+ means $r r^*$ and
 - $r?$ means $r | \epsilon$
 - $.$ means any chars
 - $a-z$ means chars between a and z

Recognition of Tokens

- Consider the language fragment:


```
if → if
then → then
else → else
relop → < | <= | = | <> | > | >=
id → letter (letter | digit)*
num → digit+ (.digit+)? (E(+|-)?digit+)?
```
- Assume lexemes are separated by white space. The regular expression for white space is ws.


```
delim → blank | tab | newline
ws → delim+
```
- The lexical analyzer does **not return a token for ws**. Rather, it finds a token following the whitespace and returns that to the parser.

有限自动机 FINITE AUTOMATA

Finite Automata

- Regular expressions = specification
- Finite automata = implementation

Finite Automata

- A mathematical model
 - state transition diagram
- Recognizer for a given language
- 5-tuple $\{Q, \Sigma, \delta, q_0, F\}$
 - Q is a finite set of states
 - Σ is a finite set of input
 - f transition function $Q \times \Sigma$
 - q_0, δ , initial and final state respectively

Finite Automata

- Transitions

$s_1 \xrightarrow{a} s_2$
- Is read


In state s_1 on input "a" go to state s_2
- If end of input and in accepting state \Rightarrow accept
- Otherwise \Rightarrow reject

Finite Automata State Graphs

- A state
- The start state
- An accepting state
- A transition


A Simple Example

- A finite automaton that accepts




Another Simple Example

- A finite automaton accepting any number of 1's followed by a single 0
- Alphabet: $\{0,1\}$




And Another Example

- Alphabet $\{0,1\}$
- What language does this recognize?



Epsilon Moves


- Another kind of transition: ϵ -moves



- Machine can move from state A to state B without reading input

Finite Automata(FA)

- NFA(Nondeterministic FA) vs. DFA(Deterministic FA)
 - Represented by a directed graph
 - NFA: But different rule applications may yield different final results
 - The same $f(s, i)$ results in a different state
- DFA is a special case of NFA
 - No state has an ϵ transition
 - For each state s and input a , there is at most one edge labeled a leaving s .
 - Give examples (see the board)
- Conversion NFA \rightarrow DFA



Transition Diagrams

复旦大学媒体计算研究所

- Double circles mark *accepting states*; where a token has been found.
- Asterisks marks states where a character must be pushed back.
- E.g. Identifiers and keywords

复旦大学媒体计算研究所

Keywords

- 可以采用以下两种方式
- (1) write a separate transition diagram for each keyword
- (2) load the keywords in the symbol table before reading source (a field in the symbol table entry contains the token for the keyword, for non-keywords the field contains the id token).

复旦大学媒体计算研究所

Outline

- Specifying lexical structure using regular expressions
- Finite automata
 - Deterministic Finite Automata (DFAs)
 - Non-deterministic Finite Automata (NFAs)
- Implementation of regular expressions
 - RegExp \Rightarrow NFA \Rightarrow DFA \Rightarrow Tables

复旦大学媒体计算研究所

Regular Expressions to Finite Automata

- High-level sketch

复旦大学媒体计算研究所

FA \rightarrow RegExp

a. $1 \xrightarrow{R1} 2 \xrightarrow{R2} 3 \Rightarrow 1 \xrightarrow{R1R2} 3$

b. $1 \xrightarrow{R1} 2 \xrightarrow{R2} 1 \Rightarrow 1 \xrightarrow{R1|R2} 2$

c. $1 \xrightarrow{R1} 2 \xrightarrow{R2} 2 \xrightarrow{R3} 3 \Rightarrow 1 \xrightarrow{R1R2^*R3} 3$


复旦大学媒体计算研究所

FA \rightarrow RegExp

$L(R) = (a|b)(a|b)^*$

Regular Expressions to NFA (1)

- For ϵ ,



- For a ,



Thompson's Construction

Regular Expressions to NFA (2)

- For $s | t$,



Thompson's Construction

Regular Expressions to NFA (3)


- For st ,



Thompson's Construction

Regular Expressions to NFA (4)

- For s^* ,



Thompson's Construction

Thompson's Construction


- $N(r)$ has at most twice as many states as the number of symbols and operators in r .
- $N(r)$ has exactly one start state and one accepting state.
- Each state of $N(r)$ has either one outgoing transition on a symbol in Σ or at most two outgoing ϵ -transitions.

Example of RegExp \rightarrow NFA conversion

- $(a|b)^*abb$

Acceptance of NFAs

- Rule: NFA accepts if it can get to a final state
- Weakness
 - An NFA can get into multiple states
- Example



How to judge?

- Backtracking
- Save all possible reaching states

ϵ -closure

- ϵ -closure(s)
 - Set of NFA states reachable from NFA state s on ϵ -transitions along.
- ϵ -closure(T)
 - Set of NFA states reachable from NFA state s in T on ϵ -transitions along.
- move(T, a)
 - Set of NFA states to which there is a transition on input symbol a from some NFA state s in T

Simulating the NFA

```

S =  $\epsilon$ -closure({s0});
a = nextchar;
while (a != eof)
{
    S =  $\epsilon$ -closure(move(S, a));
    a = nextchar;
}
if (S  $\cap$  F !=  $\emptyset$ )
    return 1;
else
    return 0;

```

NFA to DFA

- NFAs and DFAs recognize the same set of languages (regular languages)
- DFAs are faster to execute
 - There are no choices to consider
- For a given language NFA can be simpler than DFA

NFA to DFA: The Trick

- Simulate the NFA
- Each state of DFA
 - = a non-empty subset of states of the NFA
- Start state
 - = the set of NFA states reachable through ϵ -moves from NFA start state
- Add a transition $S \xrightarrow{a} S'$ to DFA iff
 - S' is the set of NFA states reachable from any state in S after seeing the input a , considering ϵ -moves as well

Subset Construction

- NFA: N , DFA: D ,
- Construct $Dstates$ for D
- Construct a transition table $Dtran$ for D .
- Algorithm:
 - Initially, ϵ -closure(s_0) is only state in $Dstates$ and it is unmarked

算法

```

while (there is an unmarked state T in
Dstates) {
  Mark T;
  for each input symbol a
  {
    U =  $\epsilon$ -closure(move(T,a));
    if (U is not in Dstates)
      Add U as an unmarked state to Dstates;
    Dtran[T,a]=U;
  }
}

```

NFA to DFA. Remark

- An NFA may be in many states at any time
- How many different states?
- If there are N states, the NFA must be in some subset of those N states
- How many subsets are there?
 - $2^N - 1$ = finitely many

Implementation

- A DFA can be implemented by a 2D table T
 - One dimension is “states”
 - Other dimension is “input symbol”
- For every transition $S_i \xrightarrow{a} S_k$ define $T[i,a] = k$
- DFA “execution”
 - If in state S_i and input a , read $T[i,a] = k$ and skip to state S_k
 - Very efficient

Implementation (Cont.)

- NFA \rightarrow DFA conversion is at the heart of tools, such as flex
- But, DFAs can be huge
- In practice, flex-like tools trade off speed for space in the choice of NFA and DFA representations

Automaton	Space	Time
NFA	$O(r)$	$O(x \cdot r)$
DFA	$O(2^{ r })$	$O(x)$

$|r|$ 正则表达式长度, $|x|$ 表示输入字符串的长度

DFA minimization

- DFA minimization is the task of transforming a given deterministic finite automaton (DFA) into an equivalent DFA that has minimum number of states.
- Here, two DFAs are called equivalent if they describe the same regular language.
- Several different algorithms accomplishing this task are known and described in standard textbooks on automata theory.

DFA minimization

- For each regular language that can be accepted by a DFA,
 - there exists a DFA with a minimum number of states (and thus a minimum programming effort to create and use)
 - and this DFA is unique (except that states can be given different names.)

Minimum DFA


- There are three classes of states that can be removed/merged from the original DFA without affecting the language it accepts.
- Unreachable states** are those states that are not reachable from the initial state of the DFA, for any input string.
- Dead states** are those non-accepting states whose transitions for every input character terminate on themselves. These are also called Trap states because once entered there is no escape.
- Non-distinguishable states** are those that cannot be distinguished from one another for any input string.

Minimize DFA state number

- Merging
- Split

Merging

- Two states **s1** and **s2** are equivalent when the machine starting in s1 accepts a string δ if and only if starting in s2 it accepts δ .
- How to find equivalent states?
 - $\text{trans}[s1, c] = \text{trans}[s2, c]$
 - not sufficiently general




Split

- There are two major groups: **non-final** and **final** states.
- A group is split on a terminal if the states in the group have transitions on this terminal to different groups.
- If all the transitions are to the same group, then there is no split on this group based on this terminal.


```
do{
  for each group G,
    partition G if G contains distinguishable states.
} while( the number of groups doesn't increase.)
```

Use FA to build scanner




Example :

$a \mid abb \mid a^*b^+$



- Max length match
 - Multi-match select
- Input: aaba; abb



scanner comparison

- Hand-coded scanner (like any ordinary program):
 - Programmer creates types, defines data & procedures, designs flow of control, implements in source language.
- Lex-generated scanner:
 - Programmer writes patterns
 - (Declarative, not procedural)
 - Lex implements flow of control
 - Must less hand-coding, but
 - code looks pretty alien, tricky to debugs

Lexical Analyzer Generator


- Lex, A Lexical Analyzer Generator
 - <http://dinosaur.compilertools.net/>
- Flex, A fast scanner generator
 - <http://flex.sourceforge.net/>
- JFLEX, The Fast Scanner Generator for Java
 - <http://jflex.de/>

Testing Lexical Analyzer

- Create a suite of test source files to run through your analyzer rather than entering the source through the keyboard.
- Much faster
- More thorough
- Repeatable : You can make sure that correcting one bug in your analyzer doesn't introduce other bugs.
- Better documentation.

LEX

- A tool to generate a lexical analyzer from regular expressions.



what does lex do?

- Input: patterns written by programmer, describing tokens in language
- Process:
 - Reads patterns as regular expressions
 - Builds finite automaton to accept valid tokens
- Output: C code implementation of FA
- Compile and link C code, you've got a scanner

Structure of a Lex file

- The structure of a Lex file is divided up into three sections.

```
Definition section
%%
Rules section
%%
C code section
```


Regular Definitions

```
%%
delim      [\t\n]
ws         {delim}+
letter     [A-Za-z]
digit      [0-9]
id          {letter}({letter} | {digit})*
number     {digit} + (\.{digit} + ) ? (E [+ \-] / {digit} + ) ?
%%
{ws}       /*no action and no return*/
if         {return (IF);}
then       {return (THEN);}
else       {return (ELSE);}
{id}       {yyval = install_id(); return(ID);}
{number}   {yyval = install_num(); return(NUM);}
```

```
"<"       {yyval = LT; return(RELOP);}
"<="      {yyval = LE; return(RELOP);}
"="        {yyval = EQ; return(RELOP);}
"< >"      {yyval = NE; return(RELOP);}
">"        {yyval = GT; return(RELOP);}
">="      {yyval = GE; return(RELOP);}
%%
install_id() /*procedure to install a lexeme into the symbol table and
return a pointer thereto*/
install_num() /*procedure to install a lexeme into the number table
and return a pointer thereto*/
```

JLEX

- A tool to generate a lexical analyzer from regular expressions.
- based upon the Lex lexical analyzer generator model. JLex takes a specification file similar to that accepted by Lex, then creates a Java source file for the corresponding lexical analyzer




Sample of FLEX

```
/* Definition section */
%{
/* C code to be copied verbatim */
#include <stdio.h>
%}
/* This tells flex to read only one input file */
%option noyywrap

/* Rules section */
/* [0-9]+ matches a string of one or more digits */
[0-9]+ {
    /* yytext is a string containing the matched text. */
    printf("Saw an integer: %s\n", yytext);
}
.\n { /* Ignore all other characters. */ }
```


```
%%
/* C Code section */

int main(void)
{
    /* Call the lexer, then quit. */
    yylex();
    return 0;
}
```




复旦大学
媒体计算
研究所

Summary:




- Automata theory \Rightarrow Programming practice
 - Regular expressions and automata theory prove you can write regular expressions, give them to a program like `lex`, which will generate a machine to accept exactly those expressions.



复旦大学
媒体计算
研究所

课后作业



- 2.3a
- 2.4b
- 2.5a
- 2.6
- 用正则表达式描述Java程序中的注释
- 用正则表达式描述C++程序中的注释