

Support Vector Machines

Mingmin Chi


SCS Fudan University, Shanghai, China

- 1 Linear Separable Support Vector Machines
 - Large Margin Classifiers
 - Solution of SVMs
- 2 Linear Non-Separable SVMs
- 3 Non-Linear SVMs
- 4 Multi-Class Classification Problems

- 1 **Linear Separable Support Vector Machines**
 - Large Margin Classifiers
 - Solution of SVMs
- 2 Linear Non-Separable SVMs
- 3 Non-Linear SVMs
- 4 Multi-Class Classification Problems

- 1 **Linear Separable Support Vector Machines**
 - **Large Margin Classifiers**
 - Solution of SVMs
- 2 Linear Non-Separable SVMs
- 3 Non-Linear SVMs
- 4 Multi-Class Classification Problems

Decision Boundary of Perceptron



Convergence of the Perceptron

- Suppose that there exists the optimal solution (\mathbf{w}^*, b^*) , which defines a decision boundary correctly classifying all the training samples, and every training sample is at least distance $\rho > 0$ from the decision boundary, i.e.,

Convergence of the Perceptron

- Suppose that there exists the optimal solution (\mathbf{w}^*, b^*) , which defines a decision boundary correctly classifying all the training samples, and every training sample is at least distance $\rho > 0$ from the decision boundary, i.e.,

$$|f(\mathbf{x}_i)| = |\mathbf{w}^{*\top} \mathbf{x}_i + b^*| = \rho$$

- Suppose that there exists a $\rho > 0$, and a weight vector \mathbf{w}^* satisfying $\|\mathbf{w}^*\| = 1$, and a threshold b^* , such that

Convergence of the Perceptron

- Suppose that there exists the optimal solution (\mathbf{w}^*, b^*) , which defines a decision boundary correctly classifying all the training samples, and every training sample is at least distance $\rho > 0$ from the decision boundary, i.e.,


$$|f(\mathbf{x}_i)| = |\mathbf{w}^{*\top} \mathbf{x}_i + b^*| = \rho$$

- Suppose that there exists a $\rho > 0$, and a weight vector \mathbf{w}^* satisfying $\|\mathbf{w}^*\| = 1$, and a threshold b^* , such that

$$\forall_{i=1}^n, y_i f(\mathbf{x}_i) \geq \rho$$

- Then the perceptron algorithm converges after no more than $(b^{*2} + 1)(R^2 + 1)/\rho^2$ updates, where $R = \max_i \|\mathbf{x}_i\|$. [Novikov, 1962]

Distance From Decision Boundary



Definition of Margin

- The quantity ρ plays a crucial role for the perceptron as it determines

Definition of Margin

- The quantity ρ plays a crucial role for the perceptron as it determines (1) how well the two classes can be separated;

Definition of Margin

- The quantity ρ plays a crucial role for the perceptron as it determines (1) how well the two classes can be separated; (2) How fast the perceptron learning algorithm converges
- This quantity ρ is henceforth what we call a *margin*

Definition of Margin

- The quantity ρ plays a crucial role for the perceptron as it determines (1) how well the two classes can be separated; (2) How fast the perceptron learning algorithm converges
- This quantity ρ is henceforth what we call a *margin*
- **Definition:** Denote by $f : \mathbb{R}^d \rightarrow \mathbb{R}$ used for the classification, then

$$\rho_f(\mathbf{x}, y) := yf(\mathbf{x})$$

Definition of Margin

- The quantity ρ plays a crucial role for the perceptron as it determines (1) how well the two classes can be separated; (2) How fast the perceptron learning algorithm converges
- This quantity ρ is henceforth what we call a *margin*
- **Definition:** Denote by $f : \mathbb{R}^d \rightarrow \mathbb{R}$ used for the classification, then

$$\rho_f(\mathbf{x}, y) := yf(\mathbf{x})$$

- Denote the minimum margin over the whole samples

$$\rho_f := \min_{1 \leq i \leq n} \rho_f(\mathbf{x}_i, y_i)$$

Maximum Margin Hyperplane

- It is desirable to have an estimator with a large margin
- Question?

Maximum Margin Hyperplane

- It is desirable to have an estimator with a large margin
- Question? Whether there exists such estimator with *maximum* margin,

Maximum Margin Hyperplane

- It is desirable to have an estimator with a large margin
- Question? Whether there exists such estimator with *maximum* margin, i.e., whether some f^* exists with

$$f^* := \arg \max_f \rho_f = \arg \max_f \min_i \rho_f(\mathbf{x}_i, y_i)$$

Maximum Margin Hyperplane

- It is desirable to have an estimator with a large margin
- Question? Whether there exists such estimator with *maximum* margin, i.e., whether some f^* exists with

$$f^* := \arg \max_f \rho_f = \arg \max_f \min_i \rho_f(\mathbf{x}_i, y_i)$$

- Without the constraints on the size of \mathbf{w} , this maximum does not exist

Maximum Margin Hyperplane (Contd)

- If we define $f : \mathbb{R}^d \rightarrow \mathbb{R}$ by

$$f(\mathbf{x}) = \frac{\mathbf{w}^\top \mathbf{x} + b}{\|\mathbf{w}\|},$$

then the maximum margin f is defined by the weight vector and threshold that satisfy

$$\mathbf{w}^*, b^* = \arg \max_{\mathbf{w}, b} \min_{i=1}^m \frac{y_i(\mathbf{w}^\top \mathbf{x}_i + b)}{\|\mathbf{w}\|} \quad (1)$$

Maximum Margin Hyperplane (Contd)

- If we define $f : \mathbb{R}^d \rightarrow \mathbb{R}$ by

$$f(\mathbf{x}) = \frac{\mathbf{w}^\top \mathbf{x} + b}{\|\mathbf{w}\|},$$

then the maximum margin f is defined by the weight vector and threshold that satisfy

$$\begin{aligned} \mathbf{w}^*, b^* &= \arg \max_{\mathbf{w}, b} \min_{i=1}^m \frac{y_i(\mathbf{w}^\top \mathbf{x}_i + b)}{\|\mathbf{w}\|} \\ &= \arg \max_{\mathbf{w}, b} \min_{i=1}^m y_i \operatorname{sgn}(\mathbf{w}^\top \mathbf{x}_i + b) \left\| \frac{\mathbf{w}^\top \mathbf{x}_i}{\|\mathbf{w}\|^2} \mathbf{w} + \frac{b}{\|\mathbf{w}\|^2} \mathbf{w} \right\| \end{aligned} \quad (1)$$

Maximum Margin Hyperplane (Cont)

- The Equ.(1) is equivalent to

$$\mathbf{w}^*, b^*, \rho^* = \arg \max_{\mathbf{w}, b, \rho} \rho \quad \text{s.t.} \quad \frac{y_i(\mathbf{w}^\top \mathbf{x}_i + b)}{\|\mathbf{w}\|} \geq \rho$$

Maximum Margin Hyperplane (Cont)

- The Equ.(1) is equivalent to

$$\begin{aligned}\mathbf{w}^*, b^*, \rho^* &= \arg \max_{\mathbf{w}, b, \rho} \rho \quad \text{s.t.} \quad \frac{y_i(\mathbf{w}^\top \mathbf{x}_i + b)}{\|\mathbf{w}\|} \geq \rho \\ &= \arg \max_{\mathbf{w}, b, \rho} \rho \quad \text{s.t.} \quad \|\mathbf{w}\| = 1 \text{ and } y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq \rho\end{aligned}$$

Maximum Margin Hyperplane (Cont)

- The Equ.(1) is equivalent to

$$\begin{aligned}\mathbf{w}^*, b^*, \rho^* &= \arg \max_{\mathbf{w}, b, \rho} \rho \quad \text{s.t.} \quad \frac{y_i(\mathbf{w}^\top \mathbf{x}_i + b)}{\|\mathbf{w}\|} \geq \rho \\ &= \arg \max_{\mathbf{w}, b, \rho} \rho \quad \text{s.t.} \quad \|\mathbf{w}\| = 1 \text{ and } y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq \rho \\ &= \arg \max_{\mathbf{w}, b} \|\mathbf{w}\|^2 \quad \text{s.t.} \quad y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1\end{aligned}$$

where $1 \leq i \leq n$

Canonical Hyperplanes

- **Definition:** The hyperplane is in *canonical* form w.r.t. \mathbf{X} if $\min_{\mathbf{x}_i \in \mathbf{X}} |\mathbf{w}^\top \mathbf{x}_i + b| = 1$

Canonical Hyperplanes


- **Definition:** The hyperplane is in *canonical* form w.r.t. \mathbf{X} if $\min_{\mathbf{x}_i \in \mathbf{X}} |\mathbf{w}^\top \mathbf{x}_i + b| = 1$
- For canonical hyperplanes, the distance of the closest point to the hyperplane ("margin") is $1/\|\mathbf{w}\|$: Assume we have two points on the two margins: i.e., \mathbf{x}^+ on the positive margin and \mathbf{x}^- on the negative margin, so we have

$$\begin{aligned}f(\mathbf{x}^+) &= \langle \frac{\mathbf{w}}{\|\mathbf{w}\|}, \mathbf{x}^+ \rangle + \frac{b}{\|\mathbf{w}\|} = \frac{1}{\|\mathbf{w}\|} \\f(\mathbf{x}^-) &= \langle \frac{\mathbf{w}}{\|\mathbf{w}\|}, \mathbf{x}^- \rangle + \frac{b}{\|\mathbf{w}\|} = -\frac{1}{\|\mathbf{w}\|}\end{aligned}$$

then the geometric margin ρ is then the functional margin of the resulting classifier

$$\rho = \frac{1}{2} (f(\mathbf{x}^+) - f(\mathbf{x}^-)) = \frac{1}{\|\mathbf{w}\|}$$

Separable (Hard-Margin) SVMs



Separable (Hard-Margin) SVMs

For the ease of computation, finally we can write down the objective function with the constraints:

$$\begin{aligned} \min_{\mathbf{w}} \quad & \left\{ \frac{1}{2} \|\mathbf{w}\|^2 \right\} \\ \text{s.t.} \quad & \forall_{i=1}^n : y_i (\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 \end{aligned}$$

This problem can be solved in the primal and dual formulations. usually dealt with by the Lagrange theory, where we can introduce Lagrange multipliers $\alpha_i \geq 0$ and a *Lagrangian* as follows:

$$L(\mathbf{w}, b, \alpha) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^n \alpha_i \left(y_i (\mathbf{w}^\top \mathbf{x}_i + b) - 1 \right)$$

- 1 Linear Separable Support Vector Machines
 - Large Margin Classifiers
 - **Solution of SVMs**
- 2 Linear Non-Separable SVMs
- 3 Non-Linear SVMs
- 4 Multi-Class Classification Problems

Dual Problems

- This problem is usually transformed to its corresponding dual form by introducing Lagrange multipliers $\alpha_i \geq 0$.
- The primal *Lagrangian* is:

$$L(\mathbf{w}, b, \alpha) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^n \alpha_i \left(y_i (\mathbf{w}^\top \mathbf{x}_i + b) - 1 \right),$$

here, we define \mathbf{w} and b as the *primal variables* and α_i as the *dual variables*

- The corresponding dual is found by differentiating with respect to the primal variables \mathbf{w} and b due to the *Karush-Kuhn-Tucker (KKT) conditions*

Dual Problems (Contd)

- According to the KKT conditions, at the optimal point, the derivatives of the Lagrangian $L(\mathbf{w}, b, \alpha)$ with respect to the primal variables must vanish
- Hence we have:

$$\frac{\partial L(\mathbf{w}, b, \alpha)}{\partial \mathbf{w}} = \mathbf{w} - \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i = 0 \Rightarrow \mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$

and
$$\frac{\partial L(\mathbf{w}, b, \alpha)}{\partial b} = 0 \Rightarrow \sum_{i=1}^n \alpha_i y_i = 0$$

Dual Lagrangian

- Substituting

$$\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i \text{ and } \sum_{i=1}^n \alpha_i y_i = 0$$

into the primal $L(\mathbf{w}, b, \alpha) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^n \alpha_i (y_i (\mathbf{w}^\top \mathbf{x}_i + b) - 1)$,

to obtain

$$L(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{l=1}^n y_i y_l \alpha_i \alpha_l \langle \mathbf{x}_i, \mathbf{x}_l \rangle$$

s.t. $\begin{cases} \alpha_i \geq 0, 1 \leq i \leq n \\ \sum_{i=1}^n y_i \alpha_i = 0 \end{cases}$

Computing b

$$f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + b = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i^\top \mathbf{x} + b$$

Since any support vector \mathbf{x}_j satisfies $y_j f(\mathbf{x}_j) = 1$, we have

$$\begin{aligned} y_j f(\mathbf{x}_j) &= y_j \left(\sum_{i \in \mathcal{S}} \alpha_i y_i \mathbf{x}_i^\top \mathbf{x}_j + b \right) = 1 \\ \Rightarrow f(\mathbf{x}_j) &= \sum_{i \in \mathcal{S}} \alpha_i y_i \mathbf{x}_i^\top \mathbf{x}_j + b = y_j \end{aligned}$$

By averaging these over all support vectors

$$b^* = \frac{1}{N_S} \sum_{j \in \mathcal{S}} \left(y_j - \sum_{i \in \mathcal{S}} \alpha_i y_i \mathbf{x}_i^\top \mathbf{x}_j \right)$$

Decision Function

$$\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$

The decision function

$$\begin{aligned} f(\mathbf{x}) &= \mathbf{w}^\top \mathbf{x} + b \\ &= \text{sgn} \left[\sum_{i=1}^n \alpha_i^* y_i \mathbf{x}_i^\top \mathbf{x} + b^* \right] \end{aligned}$$

Support Vectors

Karush-Kuhn-Tuck (KKT) conditions:

$$\alpha_i \geq 0$$


$$y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1$$

$$\alpha_i \left(y_i(\mathbf{w}^\top \mathbf{x}_i + b) - 1 \right) = 0$$

Data points are


- 1 no support vectors if $\alpha_i = 0 : y_i(\mathbf{w}^\top \mathbf{x}_i + b) > 1$
- 2 support vectors if $\alpha_i > 0 : y_i(\mathbf{w}^\top \mathbf{x}_i + b) = 1$


Support Vectors (Contd)



- 1 Linear Separable Support Vector Machines
 - Large Margin Classifiers
 - Solution of SVMs
- 2 Linear Non-Separable SVMs
- 3 Non-Linear SVMs
- 4 Multi-Class Classification Problems

Question?





Soft-Margin SVMs

- To overcome the sensitivity to the noisy data, a standard approach is to allow for the possibility of example violating the critical constraints by introducing “slack variable”:

$$\xi_i \geq 0, \text{ for all } 1 \leq i \leq n, \quad (2)$$

along with the relaxed constraints:

$$y_i \left(\mathbf{w}^\top \mathbf{x}_i + b \right) \geq 1 - \xi_i, \text{ for all } 1 \leq i \leq n \quad (3)$$

- By making ξ_i large enough, the constraint on (\mathbf{x}_i, y_i) can always be met.

Soft-Margin SVMs (Contd)

- In order not to obtain the trivial solution where all ξ_i take on large values, we thus need to penalize them in the objective function.
- To this end, a term $\sum_i \xi_i$ is included in the objective function as

$$\min_{\mathbf{w}, \xi} \left\{ \frac{1}{2} \|\mathbf{w}\|^2 + c \sum_{i=1}^n \xi_i \right\}$$

subject to (3) and (2)

Dual Problems

By introducing Lagrange multipliers $\alpha_i \geq 0$ and $\mu_i \geq 0$. The primal *Lagrangian* is:

$$L(\mathbf{w}, b, \alpha) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i - \sum_{i=1}^n \alpha_i \left(y_i (\mathbf{w}^\top \mathbf{x}_i + b) - 1 \right) - \sum_{i=1}^n \mu_i \xi_i,$$

$$\frac{\partial L}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$

$$\frac{\partial L}{\partial b} = 0 \Rightarrow \sum_{i=1}^n \alpha_i y_i = 0$$

$$\frac{\partial L}{\partial \xi_i} = 0 \Rightarrow \alpha_i = C - \xi_i$$

Solution

With the Lagrange theory and the KKT conditions, we can obtain the dual Lagrangian:

$$L(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{l=1}^n y_i y_l \alpha_i \alpha_l \mathbf{x}_i^T \mathbf{x}_l$$

$$\text{s.t. } \begin{cases} \alpha_i \geq 0, 1 \leq i \leq n \\ \sum_{i=1}^n y_i \alpha_i = 0 \end{cases}$$

- the box constraint
- the optimal value of \mathbf{w} in terms of the optimal value of α^*

$$\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$

Support Vectors

With KKT conditions:

$$\alpha_i \geq 0; \mu_i \geq 0; \xi_i \geq 0; \mu_i \xi_i = 0$$

$$y_i(\mathbf{w}^\top \mathbf{x}_i + b) - 1 + \xi_i \geq 0$$


$$\alpha_i \left(y_i(\mathbf{w}^\top \mathbf{x}_i + b) - 1 + \xi_i \right) = 0$$

and $\alpha_i = C - \mu_i$


The data points are

- ① no contribution to \mathbf{w} if $\alpha_i = 0$
- ② support vectors if $\alpha_i > 0 \Rightarrow y_i(\mathbf{w}^\top \mathbf{x}_i + b) = 1 - \xi_i$
 - $\alpha_i < C \Rightarrow \xi_i = 0$, on the margin
 - $\alpha_i = C \Rightarrow \xi_i > 0$
 - ① correctly classified if $\xi_i \leq 1$
 - ② misclassified if $\xi_i > 1$

Soft-Margin SVMs: Geometric Illustration



Soft-Margin SVMs: Geometric Illustration (cont)



- 1 Linear Separable Support Vector Machines
 - Large Margin Classifiers
 - Solution of SVMs
- 2 Linear Non-Separable SVMs
- 3 Non-Linear SVMs
- 4 Multi-Class Classification Problems

Feature Spaces

- Preprocess the data with

$$\begin{aligned}\phi : \mathbf{X} &\rightarrow \mathcal{H} \\ \mathbf{x} &\rightarrow \phi(\mathbf{x}),\end{aligned}$$



where \mathcal{H} is a dot product space and learn the mapping from $\phi(\mathbf{x})$ to the output y

- Usually, $\dim(\mathbf{X}) \ll \dim(\mathcal{H})$

Nonlinear Mapping

$$\phi: \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

$$(\mathbf{x}_1, \mathbf{x}_2) \rightarrow (\mathbf{z}_1, \mathbf{z}_2, \mathbf{z}_3) : (\mathbf{x}_1^2, \sqrt{2}\mathbf{x}_1\mathbf{x}_2, \mathbf{x}_2^2)$$



Dual Lagrangian with Kernels

- In linear separable case, the dual formulation is

$$L(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{l=1}^n y_i y_l \alpha_i \alpha_l \mathbf{x}_i^T \mathbf{x}_l$$

$$\text{s.t. } \begin{cases} \alpha_i \geq 0, 1 \leq i \leq n \\ \sum_{i=1}^n y_i \alpha_i = 0 \end{cases}$$

Dual Lagrangian with Kernels

- In linear separable case, the dual formulation is

$$L(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{l=1}^n y_i y_l \alpha_i \alpha_l k(\mathbf{x}_i, \mathbf{x}_l)$$

$$\text{s.t. } \begin{cases} \alpha_i \geq 0, 1 \leq i \leq n \\ \sum_{i=1}^n y_i \alpha_i = 0 \end{cases}$$

- The solution

$$f(\mathbf{x}) = \sum_{i=1}^n \alpha_i^* y_i \mathbf{x}^\top \mathbf{x}_i + b^*$$

can be formulated as

$$f(\mathbf{x}) = \sum_{i=1}^n \alpha_i^* y_i k(\mathbf{x}, \mathbf{x}_i) + b^*$$

Positive Definite Kernels

- **Definition (Gram Matrix):** Given a function $k : \mathcal{X}^2 \rightarrow \mathbb{K}$ (where $\mathbb{K} = \mathcal{C}$ or $\mathbb{K} = \mathbb{R}$) and patterns $(\mathbf{x}_1, \dots, \mathbf{x}_n) \in \mathcal{X}$, the $n \times n$ matrix \mathbf{K} with elements:

$$\mathbf{K}_{ij} := k(\mathbf{x}_i, \mathbf{x}_j)$$

is called the Gram matrix (or kernel matrix) of k w.r.t $(\mathbf{x}_1, \dots, \mathbf{x}_n)$

- **Definition (Positive Definite Matrix):** A complex $n \times n$ matrix \mathbf{K} satisfying

$$\sum_{ij} c_i \bar{c}_j \mathbf{K}_{ij} \geq 0$$

- 1 for all $c \in \mathcal{C}$, the matrix \mathbf{K} is called positive definite
- 2 for all $c \in \mathcal{R}$, the real symmetric $n \times n$ matrix \mathbf{K} is called positive definite

Positive Definite Kernels (Contd)

- **Definition (Positive Definite Kernel):** Let \mathcal{X} be a nonempty set. A function k on $\mathcal{X} \times \mathcal{X}$ which for all $n \in \mathbb{N}$, and all $(\mathbf{x}_1, \dots, \mathbf{x}_n) \in \mathcal{X}$ gives rise to a positive definite Gram matrix is called a positive definite (pd) kernel
- A number of different terms are used for pd kernels, such as reproducing kernel, Mercer kernel, admissible kernel, Support Vector kernel, nonnegative definite kernel, and covariance function
- Conditions of kernels:
 - 1 positivity on the diagonal of \mathbf{K} , i.e., $k(\mathbf{x}, \mathbf{x}) > 0$, for all $\mathbf{x} \in \mathcal{X}$
 - 2 Symmetry, i.e., $k(\mathbf{x}_i, \mathbf{x}_j) = \overline{k(\mathbf{x}_j, \mathbf{x}_i)}$ or $\mathbf{K}_{ij} = \overline{\mathbf{K}_{ji}}$

Mercer Theorem

- **Mercer Theorem** If k is a continuous kernel of a positive definite integral operator on $L_2(\mathcal{X})$,

$$\int_{\mathcal{X}^2} k(\mathbf{x}, \mathbf{x}') f(\mathbf{x}) f(\mathbf{x}') d\mathbf{x} d\mathbf{x}' \geq 0, \quad \forall f$$

It can be expanded as

$$k(\mathbf{x}, \mathbf{x}') = \sum_{i=1}^{\infty} \lambda_i \psi_i(\mathbf{x}) \psi_i(\mathbf{x}')$$

using eigenfunctions ψ_i and eigenvalues $\lambda_i \geq 0$

Mercer Theorem (Contd)

- In this case,

$$\phi(\mathbf{x}) := \begin{pmatrix} \sqrt{\lambda_1} \psi_1(\mathbf{x}) \\ \sqrt{\lambda_2} \psi_2(\mathbf{x}) \\ \vdots \\ \vdots \\ \vdots \end{pmatrix}$$

- **Proposition (Mercer Kernel Map)** If k is a kernel satisfying the above conditions (in the previous slide), we can construct a mapping ϕ into a space where k acts as a dot product,

$$\langle \phi(\mathbf{x}), \phi(\mathbf{x}') \rangle = \phi(\mathbf{x})^\top \phi(\mathbf{x}') = k(\mathbf{x}, \mathbf{x}')$$

Closure Properties of Kernel Functions

Constructing more complex kernels from simpler ones

Let $k_1(\mathbf{x}, \mathbf{y})$ and $k_2(\mathbf{x}, \mathbf{y})$ be the kernels functions. Then the following are all kernels:

- ① $k_1(\mathbf{x}, \mathbf{y}) + k_2(\mathbf{x}, \mathbf{y}) \Leftarrow \phi(\mathbf{x}) = (\phi_1(\mathbf{x}), \phi_2(\mathbf{x}))^\top$
- ② $\alpha k_1(\mathbf{x}, \mathbf{y}), \alpha > 0 \Leftarrow \phi(\mathbf{x}) = \sqrt{\alpha} \phi_1(\mathbf{x})$
- ③ $k_1(\mathbf{x}, \mathbf{y}) k_2(\mathbf{x}, \mathbf{y}) \Leftarrow \phi(\mathbf{x})_{ij} = \phi_i(\mathbf{x})_i \phi_j(\mathbf{x})_j$ (tensor product)
- ④ $f(\mathbf{x}) f(\mathbf{y}), \forall f \Leftarrow \phi(\mathbf{x}) = f(\mathbf{x})$
- ⑤ $\mathbf{x}^\top \mathbf{A} \mathbf{y}, \text{ for } \mathbf{A} \succeq 0 \Leftarrow \phi(\mathbf{x}) = L^\top \mathbf{x} \text{ for } \mathbf{A} = L L^\top$ (Cholesky)

Kernel Tricks

- any algorithm that only depends on dot products can benefit from the kernel trick
- can be extended to non-vectorial data
- The kernel is as a nonlinear similarity measure (examples)
- Examples of common kernels used
 - 1 Gaussian kernels: $k(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\sigma^2}\right)$
 - 2 Polynomial kernels: $k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^\top \mathbf{x}' + c)^d$

- 1 Linear Separable Support Vector Machines
 - Large Margin Classifiers
 - Solution of SVMs
- 2 Linear Non-Separable SVMs
- 3 Non-Linear SVMs
- 4 Multi-Class Classification Problems

One-vs-Rest

- To get C-class binary classifiers, it is common to construct a set of binary classifiers f^1, \dots, f^C , each trained to separate one class from the rest, and combine them by doing the multi-class classification according to the maximal output before applying the sgn function; i.e., by taking:

$$\arg \max_{j=1, \dots, C} f^j(\mathbf{x}), \text{ where } f^j(\mathbf{x}) = \sum_{i=1}^n y_i \alpha^j k(\mathbf{x}, \mathbf{x}_i) + b^j$$

One-vs-Rest: Objective Function

- Multi-class objective functions

$$\min_{\mathbf{w}^j, \xi^j} \left\{ \frac{1}{2} \|\mathbf{w}^j\|^2 + C \sum_{i=1}^n \xi_i^j \right\}$$

$$\text{s.t.} \quad \begin{cases} \langle \mathbf{w}^j, \mathbf{x}_i \rangle + b \geq 1 - \xi_i^j, & y_i = j \\ \langle \mathbf{w}^j, \mathbf{x}_i \rangle + b \leq -1 + \xi_i^j, & y_i \neq j \\ \forall_{i=1}^n : \xi_i^j > 0, j = 1, \dots, C \end{cases}$$

One-vs-One

- This method constructs $C \times (C - 1)/2$ classifiers where each one is trained on data from two classes. For training data from the l -th and the j -th classes, we solve the following binary classification problem:

$$\begin{aligned} \min_{\mathbf{w}^{lj}, \xi^{lj}} & \left\{ \frac{1}{2} \|\mathbf{w}^{lj}\|^2 + C \sum_{i=1}^n \xi_i^{lj} \right\} \\ \text{s.t.} & \begin{cases} \langle \mathbf{w}^{lj}, \mathbf{x}_i \rangle + b \geq 1 - \xi_i^{lj}, & y_i = l \\ \langle \mathbf{w}^{lj}, \mathbf{x}_i \rangle + b \leq -1 + \xi_i^{lj}, & y_i = j \\ \forall_{i=1}^n : \xi_i^{lj} > 0, l, j = 1, \dots, C \end{cases} \end{aligned}$$

One-vs-One

- Assume the voting strategy used
 - 1 if $\text{sgn} \left[\sum_{i=1}^j \alpha_i^{lj} y_i < \mathbf{x}, \xi > + b^{lj} \right]$ says that the pattern \mathbf{x} belongs to the class j , then the vote for the class j is added by one. Otherwise the class l is added by one
 - 2 Then we predict the \mathbf{x} is in the class with the largest vote